

# Multiple Linear Regression

Oscar A. Trevizo

2023-05-07

## Contents

<b>1</b>	<b>Load the libraries</b>	<b>2</b>
<b>2</b>	<b>Multiple Linear Regression, Model Coefficients, Measures.</b>	<b>2</b>
<b>3</b>	<b>Load the data: OJ{ISLR}</b>	<b>2</b>
<b>4</b>	<b>Train Test data split</b>	<b>4</b>
<b>5</b>	<b>Example 1: Combo numeric and factors</b>	<b>4</b>
5.1	Fit the model: Train set . . . . .	4
<b>6</b>	<b>How to get model information</b>	<b>8</b>
6.1	How to get the attributes . . . . .	8
6.2	Pull the coefficients . . . . .	9
6.3	How to get more info: Pull the attributes from summary model . . . . .	9
6.4	Get the degrees of freedom: Regression, Residual Error . . . . .	9
6.5	Get the model Std. Errors . . . . .	10
6.6	Get the Confidence Intervals . . . . .	10
6.7	Get more information on the C.I. . . . .	10
6.8	t-distribution 95% level: the formula . . . . .	11
<b>7</b>	<b>Example 2: Without Intercept</b>	<b>11</b>
7.1	Fit the model without Intercept . . . . .	11
<b>8</b>	<b>Example 3: Numeric explanatory variables (3 of them)</b>	<b>14</b>
8.1	Fit the model: Train set . . . . .	14
<b>9</b>	<b>Example 4: Without Intercept</b>	<b>18</b>
9.1	Fit the model without Intercept: Train set . . . . .	18

## 1 Load the libraries

## 2 Multiple Linear Regression, Model Coefficients, Measures.

## 3 Load the data: OJ{ISLR}

Using dataset OJ from the ISLR library (see references).

```
str(OJ)
```

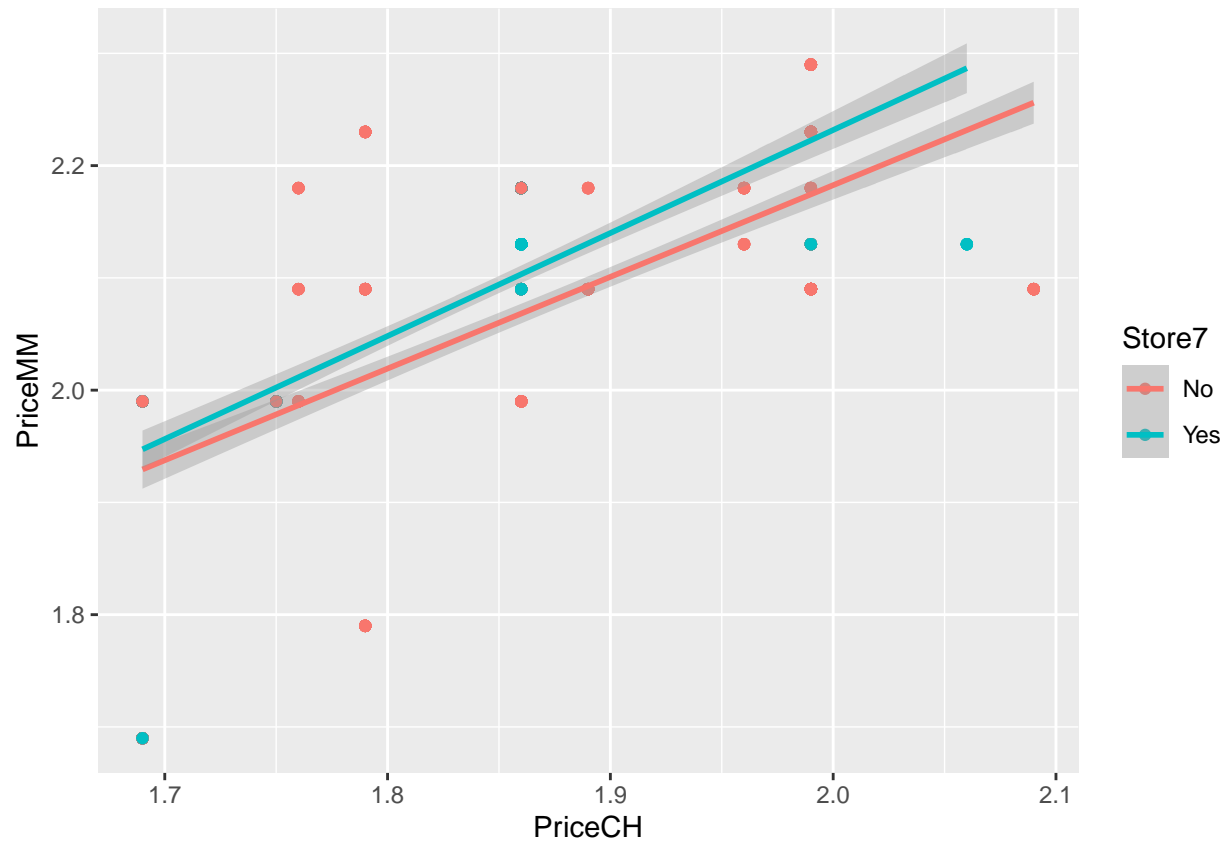
```
## 'data.frame': 1070 obs. of 18 variables:
## $ Purchase : Factor w/ 2 levels "CH","MM": 1 1 1 2 1 1 1 1 1 1 ...
## $ WeekofPurchase: num 237 239 245 227 228 230 232 234 235 238 ...
## $ StoreID : num 1 1 1 1 7 7 7 7 7 7 ...
## $ PriceCH : num 1.75 1.75 1.86 1.69 1.69 1.69 1.69 1.75 1.75 1.75 ...
## $ PriceMM : num 1.99 1.99 2.09 1.69 1.69 1.99 1.99 1.99 1.99 1.99 ...
## $ DiscCH : num 0 0 0.17 0 0 0 0 0 0 0 ...
## $ DiscMM : num 0 0.3 0 0 0 0 0.4 0.4 0.4 0.4 ...
## $ SpecialCH : num 0 0 0 0 0 0 1 1 0 0 ...
## $ SpecialMM : num 0 1 0 0 0 1 1 0 0 0 ...
## $ LoyalCH : num 0.5 0.6 0.68 0.4 0.957 ...
## $ SalePriceMM : num 1.99 1.69 2.09 1.69 1.69 1.99 1.59 1.59 1.59 1.59 ...
## $ SalePriceCH : num 1.75 1.75 1.69 1.69 1.69 1.69 1.69 1.75 1.75 1.75 ...
## $ PriceDiff : num 0.24 -0.06 0.4 0 0 0.3 -0.1 -0.16 -0.16 -0.16 ...
## $ Store7 : Factor w/ 2 levels "No","Yes": 1 1 1 1 2 2 2 2 2 2 ...
## $ PctDiscMM : num 0 0.151 0 0 0 ...
## $ PctDiscCH : num 0 0 0.0914 0 0 ...
## $ ListPriceDiff : num 0.24 0.24 0.23 0 0 0.3 0.3 0.24 0.24 0.24 ...
## $ STORE : num 1 1 1 1 0 0 0 0 0 0 ...
```

Note: There are 1070 observations. This information will be needed when calculating the degrees of freedom downstream.

### 3.0.1 ggplot PriceCH vs. PriceMM

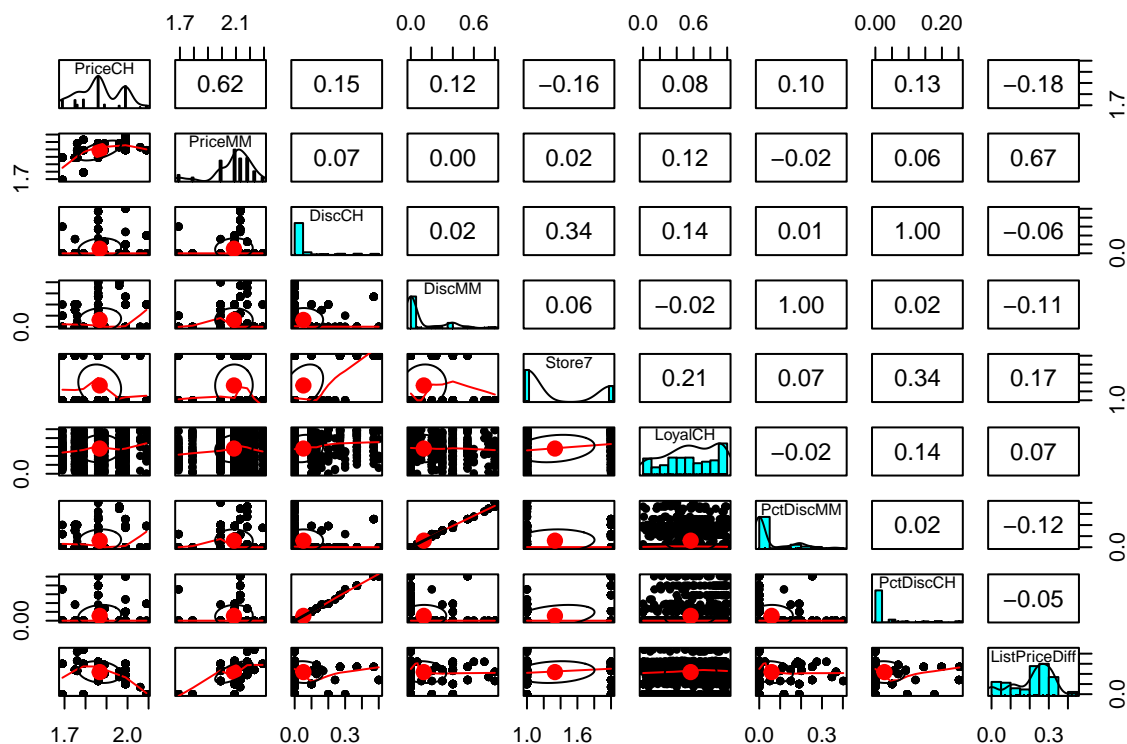
```
OJ %>% ggplot(aes(x=PriceCH, y=PriceMM, col = Store7)) +
  geom_point() +
  geom_smooth(method = 'lm')
```

```
## 'geom_smooth()' using formula = 'y ~ x'
```



### 3.0.2 pairs.panels

```
OJ %>% dplyr::select(PriceCH, PriceMM, DiscCH, DiscMM,
                    Store7, LoyalCH, PctDiscMM, PctDiscCH, ListPriceDiff) %>%
  pairs.panels()
```



## 4 Train Test data split

```
## Split the data: train / test datasets
```

```
set.seed(1234)
ind <- sample(2, nrow(OJ), replace = T, prob = c(0.7, 0.3))
train <- OJ[ind == 1,]
test <- OJ[ind == 2,]
```

```
dim(train)
```

```
## [1] 747 18
```

```
dim(test)
```

```
## [1] 323 18
```

## 5 Example 1: Combo numeric and factors

### 5.1 Fit the model: Train set

Iterate to fit the model. See *anova()* description a few chunks below.

```

# You can start with a '+' on the first variable if you want to.
# I like to start with a '+' because it can flexibly add and delete variables
# as I go through the analysis process.

```

```

m <- lm(PriceCH ~
        +PriceMM
        +WeekofPurchase
        +DiscCH
        +Store7
        +STORE
        , train)

```

```

# Display the summary with the correlation of the coefficients.
# Put it in a variable to use it downstream.

```

```

(summary_m <- summary(m, corr = TRUE))

```

```

##
## Call:
## lm(formula = PriceCH ~ +PriceMM + WeekofPurchase + DiscCH + Store7 +
##     STORE, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.129886 -0.036502 -0.006031  0.033945  0.120855
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   0.4647322   0.0337556   13.768 < 2e-16 ***
## PriceMM       0.1367351   0.0174425    7.839 1.58e-14 ***
## WeekofPurchase 0.0039965   0.0001581   25.283 < 2e-16 ***
## DiscCH       -0.0247224   0.0179649   -1.376  0.169
## Store7Yes     0.0721424   0.0068960   10.461 < 2e-16 ***
## STORE        0.0477692   0.0022143   21.573 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.04929 on 741 degrees of freedom
## Multiple R-squared:  0.7626, Adjusted R-squared:  0.761
## F-statistic: 476.1 on 5 and 741 DF, p-value: < 2.2e-16
##
## Correlation of Coefficients:
##              (Intercept) PriceMM WeekofPurchase DiscCH Store7Yes
## PriceMM          -0.32
## WeekofPurchase -0.55          -0.61
## DiscCH           0.26          0.22   -0.42
## Store7Yes       -0.07         -0.23    0.15        -0.28
## STORE           -0.05         -0.25    0.13        -0.10    0.81

```

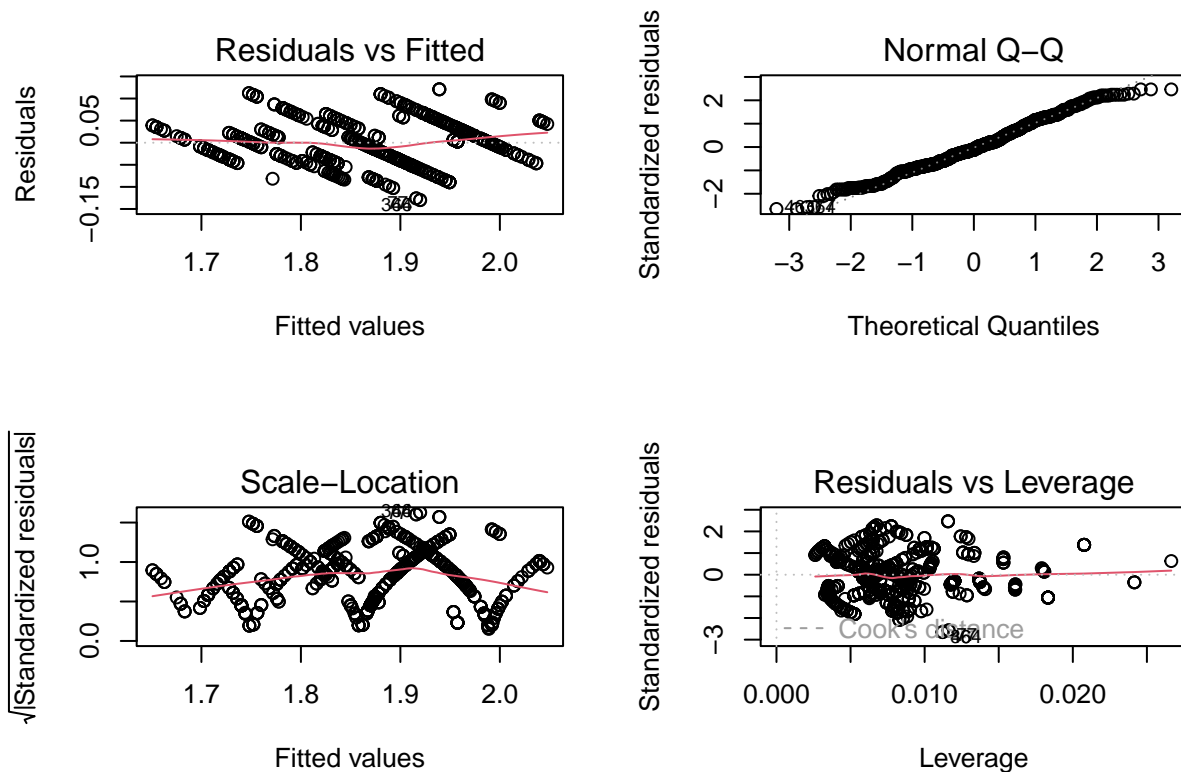
Note: Store7 is a factor variable, and it is quantified only once: Store7Yes. It is quantified once and not twice because Store7No would simply be the opposite to Store7Yes. We only need one dummy variable, therefore.

Note: Residual error has 741 degrees of freedom. That is  $= n - p - 1$ , when using the *Intercept* as we did above.

- Observations in the train dataset:  $n = 747$
- Explanatory variables:  $p = 5$
- Residual error degrees of freedom:  $DF = n - p - 1 = 747 - 5 - 1 = 741$

### 5.1.1 Diagnostic plots

```
par(mfrow = c(2,2))
plot(m)
```



```
par(mfrow = c(1,1))
```

### 5.1.2 Test the model with anova()

Compare the  $\Pr(>F)$  with the p-values from the summary of the model. They should be comparable.

```
anova(m)
```

```
## Analysis of Variance Table
##
## Response: PriceCH
##          Df Sum Sq Mean Sq F value    Pr(>F)
## PriceMM    1  2.86102   2.86102  1177.841 < 2.2e-16 ***
```

```
## WeekofPurchase    1 1.42685 1.42685 587.417 < 2.2e-16 ***
## DiscCH            1 0.02749 0.02749  11.319 0.0008067 ***
## Store7            1 0.33601 0.33601 138.329 < 2.2e-16 ***
## STORE             1 1.13047 1.13047 465.401 < 2.2e-16 ***
## Residuals         741 1.79991 0.00243
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

### 5.1.3 Predict on test dataset

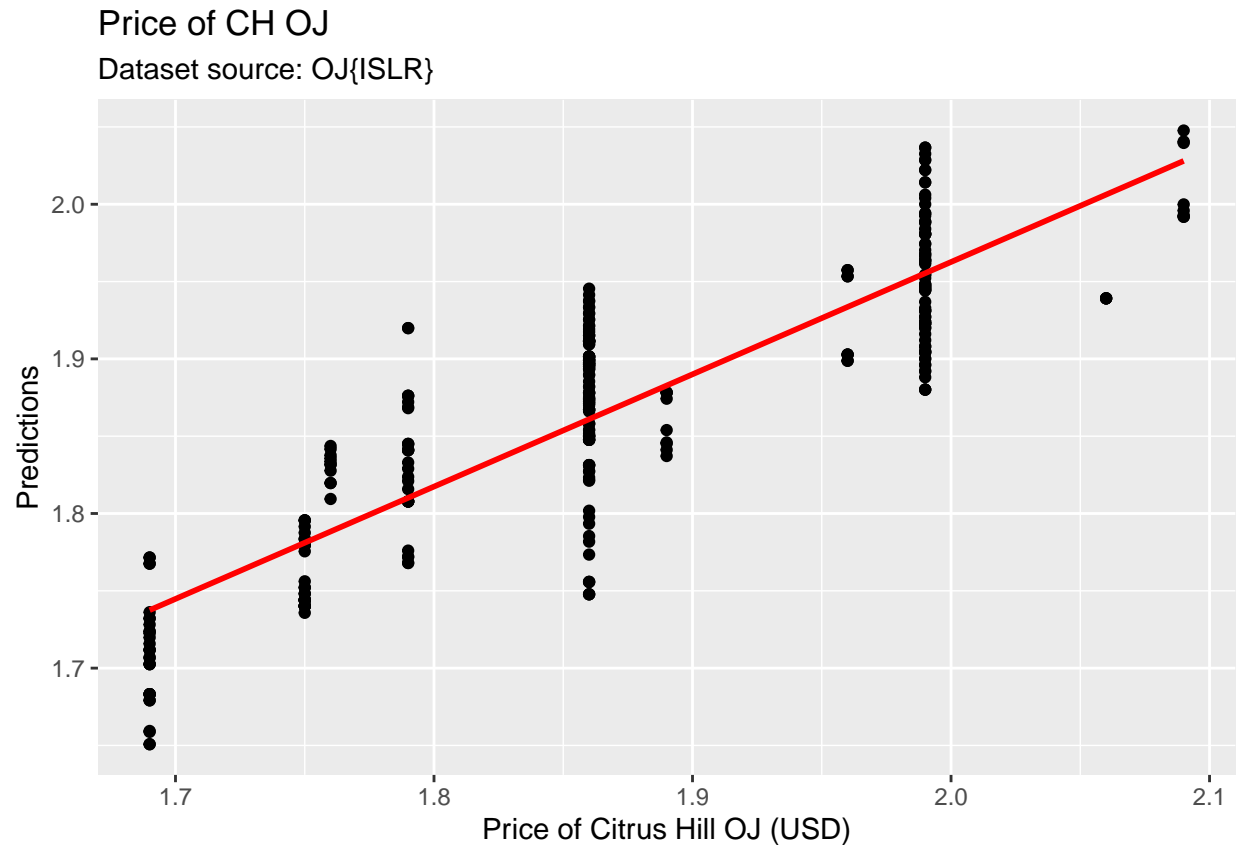
```
p <- predict(m, test)

# For ggplot we need a dataframe:
df <- data.frame(p, test)
```

### 5.1.4 Plot predictions vs actuals

```
df %>% ggplot(aes(x = PriceCH, y = p)) +
  geom_point() +
  geom_smooth(method = 'lm', col = 'red', se=FALSE) +
  scale_y_continuous('Predictions') +
  scale_x_continuous('Price of Citrus Hill OJ (USD)') +
  ggtitle('Price of CH OJ', 'Dataset source: OJ{ISLR}')

## 'geom_smooth()' using formula = 'y ~ x'
```



### 5.1.5 Assess performance: RMSE and $R^2$

- Root Mean Squared Error
- R-squared

```
# RMSE
sqrt(mean((test$PriceCH - p)^2))
```

```
## [1] 0.05199303
```

```
# R squared
cor(test$PriceCH, p)^2 ## R-Squared
```

```
## [1] 0.7539842
```

## 6 How to get model information

### 6.1 How to get the attributes

```
attributes(m)
```



```
## $names
## [1] "coefficients" "residuals" "effects" "rank"
## [5] "fitted.values" "assign" "qr" "df.residual"
## [9] "contrasts" "xlevels" "call" "terms"
## [13] "model"
##
## $class
## [1] "lm"
```

## 6.2 Pull the coefficients

```
m$coefficients
```

```
## (Intercept) PriceMM WeekofPurchase DiscCH Store7Yes
## 0.464732187 0.136735099 0.003996492 -0.024722414 0.072142374
## STORE
## 0.047769212
```

## 6.3 How to get more info: Pull the attributes from summary model

```
attributes(summary_m)
```

```
## $names
## [1] "call" "terms" "residuals" "coefficients"
## [5] "aliased" "sigma" "df" "r.squared"
## [9] "adj.r.squared" "fstatistic" "cov.unscaled" "correlation"
## [13] "symbolic.cor"
##
## $class
## [1] "summary.lm"
```

## 6.4 Get the degrees of freedom: Regression, Residual Error

```
summary_m$df
```

```
## [1] 6 741 6
```

It's giving us the degrees of freedom with  $p + theintercept = 6$ .

Worth repeating:

Regression degrees of freedom for the Residual error based on train dataset  $= n - p - 1 = 747 - 5 - 1 = 741$

[Reference DAAG page 171]

## 6.5 Get the model Std. Errors

Statology: “And to only extract the standard errors for each of the individual regression coefficients, we can use the following syntax:”

extract standard error of individual regression coefficients

```
sqrt(diag(vcov(model)))
```

```
# From: https://www.statology.org/extract-standard-error-from-lm-in-r/
# extract standard error of individual regression coefficients
# sqrt(diag(vcov(model)))
```

```
sqrt(diag(vcov(m)))
```

```
##      (Intercept)      PriceMM WeekofPurchase      DiscCH      Store7Yes
## 0.0337555941    0.0174424980    0.0001580674    0.0179648907    0.0068960344
##           STORE
## 0.0022142893
```

## 6.6 Get the Confidence Intervals

```
# Put it in a variable to get more information later on.
ci_m <- confint(m)
ci_m
```

```
##              2.5 %      97.5 %
## (Intercept)  0.398464198 0.531000176
## PriceMM      0.102492501 0.170977698
## WeekofPurchase 0.003686178 0.004306805
## DiscCH       -0.059990559 0.010545731
## Store7Yes     0.058604283 0.085680466
## STORE        0.043422185 0.052116240
```

Good to see that the confidence intervals did not cross *zero*.

## 6.7 Get more information on the C.I.

```
attributes(ci_m)
```

```
## $dim
## [1] 6 2
##
## $dimnames
## $dimnames[[1]]
## [1] "(Intercept)" "PriceMM"      "WeekofPurchase" "DiscCH"
## [5] "Store7Yes"    "STORE"
##
## $dimnames[[2]]
## [1] "2.5 %" "97.5 %"
```

95% confidence interval for volume:

$0.708 \pm 2.18 \times 0.0611$

Where:

- 0.708 = the average between the two two CI numbers (lower , higher)
- And notice how the coefficient is also the average between the CI,
- 0.708 = Volume coefficient
- 2.18 = t-value for 12 DF from qt(0.975, 12)
- 0.0611 = Standard error found from sqrt(diag(vcov(m)))[2]

### 6.7.1 Get the C.I. mean for the 2nd variable: PriceMM

```
ci_mu <- mean(c(ci_m[2,1], ci_m[2,2]))
ci_mu
```

```
## [1] 0.1367351
```

## 6.8 t-distribution 95% level: the formula

Reference DAAG page 171.

```
qt(0.975, 12)
```

```
## [1] 2.178813
```

## 7 Example 2: Without Intercept

### 7.1 Fit the model without Intercept

```
# Remove the Intercept by asdding a -1
m_noi <- lm(PriceCH ~ -1
            +PriceMM
            +WeekofPurchase
            +DiscCH
            +Store7
            +STORE
            , train)

summary(m_noi)
```

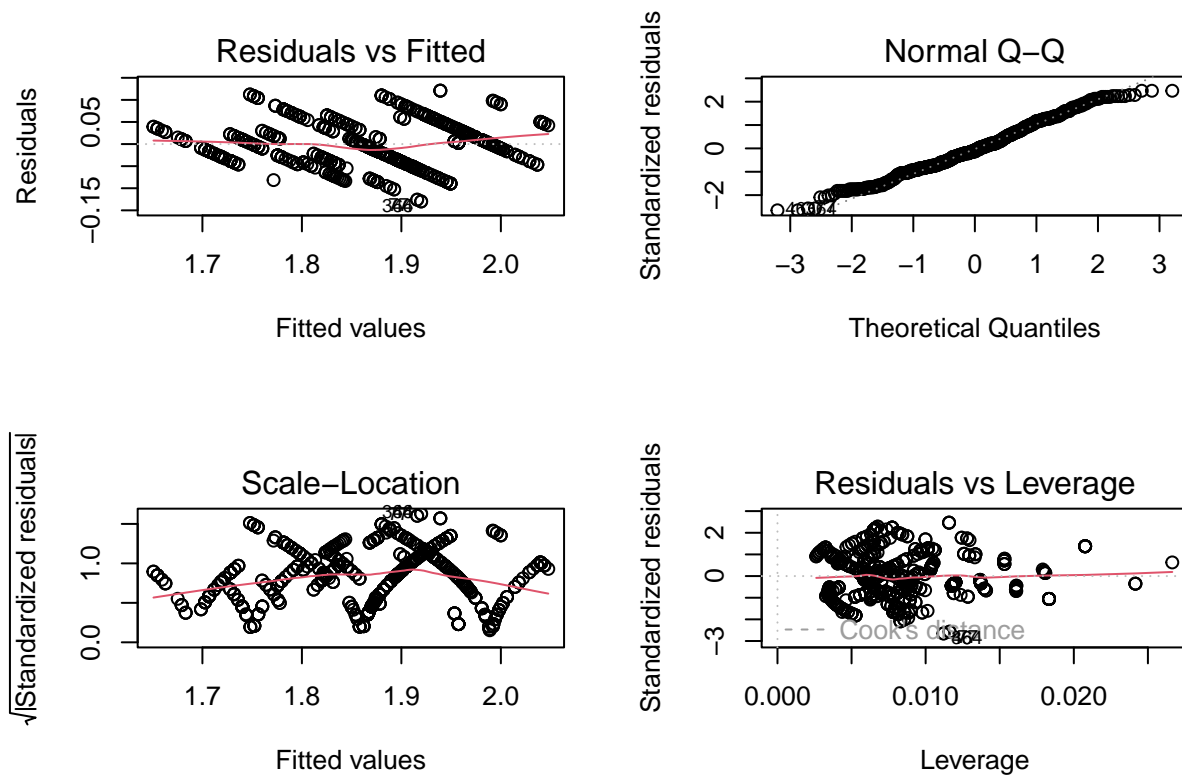
```
##
## Call:
## lm(formula = PriceCH ~ -1 + PriceMM + WeekofPurchase + DiscCH +
##      Store7 + STORE, data = train)
##
## Residuals:
```

```
##           Min           1Q       Median           3Q           Max
## -0.129886 -0.036502 -0.006031  0.033945  0.120855
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## PriceMM          0.1367351  0.0174425   7.839 1.58e-14 ***
## WeekofPurchase    0.0039965  0.0001581  25.283 < 2e-16 ***
## DiscCH            -0.0247224  0.0179649  -1.376   0.169
## Store7No           0.4647322  0.0337556  13.768 < 2e-16 ***
## Store7Yes          0.5368746  0.0339465  15.815 < 2e-16 ***
## STORE              0.0477692  0.0022143  21.573 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.04929 on 741 degrees of freedom
## Multiple R-squared:  0.9993, Adjusted R-squared:  0.9993
## F-statistic: 1.792e+05 on 6 and 741 DF, p-value: < 2.2e-16
```

Now, here pay close attention to what the model did with the factor variable.  
It created a two dummy variable, one for each level: Store7No and Store7Yes.

### 7.1.1 Plots

```
par(mfrow = c(2,2))
plot(m_noi)
```



```
par(mfrow = c(1,1))
```

### 7.1.2 Predict on test dataset

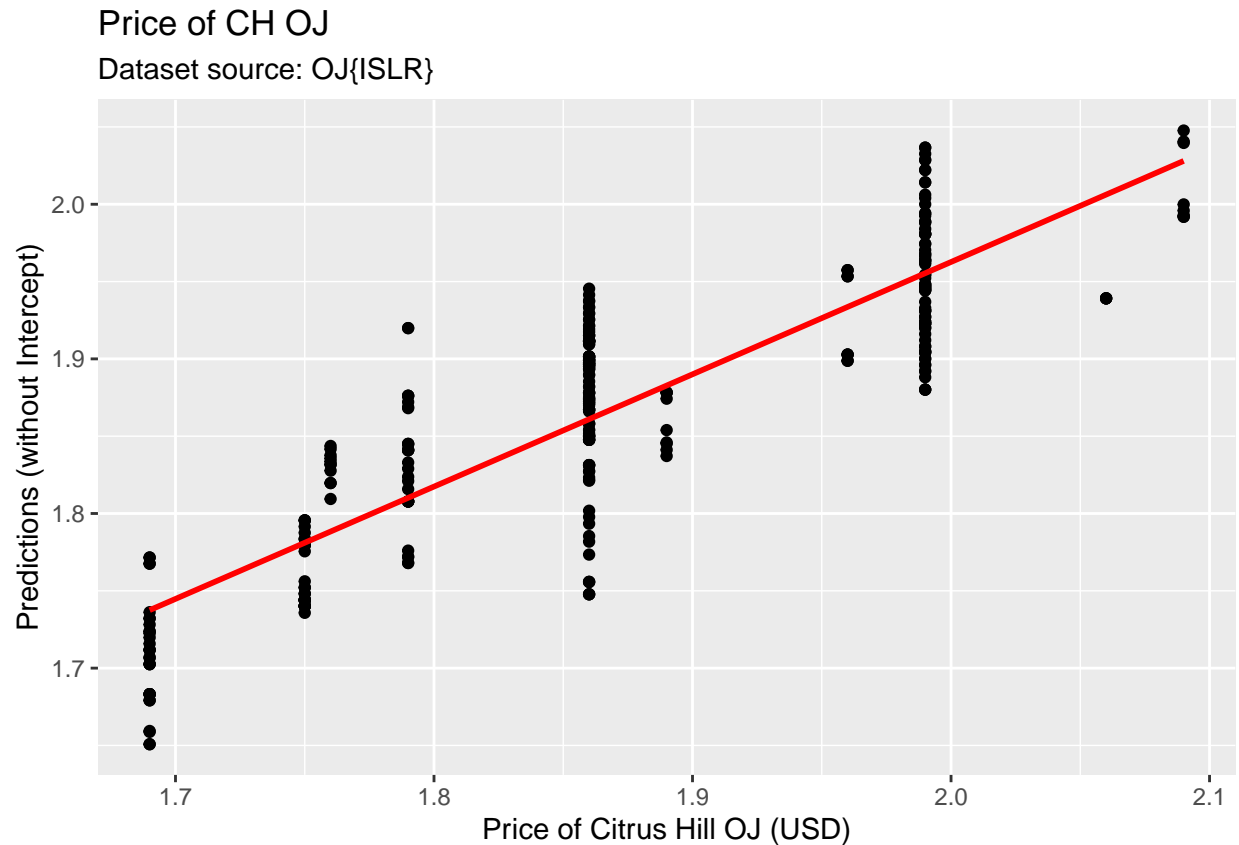
```
p_noi <- predict(m_noi, test)

# For ggplot we need a dataframe:
df_noi <- data.frame(p_noi, test)
```

### 7.1.3 Plot predictions vs actuals

```
df_noi %>% ggplot(aes(x = PriceCH, y = p_noi)) +
  geom_point() +
  geom_smooth(method = 'lm', col = 'red', se=FALSE) +
  scale_y_continuous('Predictions (without Intercept)') +
  scale_x_continuous('Price of Citrus Hill OJ (USD)') +
  ggtitle('Price of CH OJ', 'Dataset source: OJ{ISLR}')
```

```
## 'geom_smooth()' using formula = 'y ~ x'
```



#### 7.1.4 Assess performance: RMSE and $R^2$

- Root Mean Squared Error
- R-squared

```
# RMSE
sqrt(mean((test$PriceCH - p_noi)^2))
```

```
## [1] 0.05199303
```

```
# R squared
cor(test$PriceCH, p_noi)^2 ## R-Squared
```

```
## [1] 0.7539842
```

Interesting same results with or without Intercept.

## 8 Example 3: Numeric explanatory variables (3 of them)

### 8.1 Fit the model: Train set

Iterate to fit the model. See *anova()* description a few chunks below.

```

# You can start with a '+' on the first variable if you want to.
# I like to start with a '+' because it can flexibly add and delete variables
# as I go through the analysis process.
m <- lm(PriceCH ~
        +PriceMM
        +WeekofPurchase
        +STORE
        , train)

# Display the summary with the correlation of the coefficients.
# Put it in a variable to use it downstream.
(summary_m <- summary(m, corr = TRUE))

```

```

##
## Call:
## lm(formula = PriceCH ~ +PriceMM + WeekofPurchase + STORE, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.119608 -0.034881  0.005149  0.034250  0.140745
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   0.4773809   0.0349423   13.662  <2e-16 ***
## PriceMM       0.1745687   0.0179144    9.745  <2e-16 ***
## WeekofPurchase 0.0038491   0.0001538   25.021  <2e-16 ***
## STORE         0.0286057   0.0013668   20.929  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.05281 on 743 degrees of freedom
## Multiple R-squared:  0.7267, Adjusted R-squared:  0.7256
## F-statistic: 658.5 on 3 and 743 DF, p-value: < 2.2e-16
##
## Correlation of Coefficients:
##              (Intercept) PriceMM WeekofPurchase
## PriceMM          -0.40
## WeekofPurchase -0.50          -0.59
## STORE            -0.04          -0.15    0.12

```

Note: Degrees of freedom =  $n - p - 1$ , when using the *Intercept* as we did above.

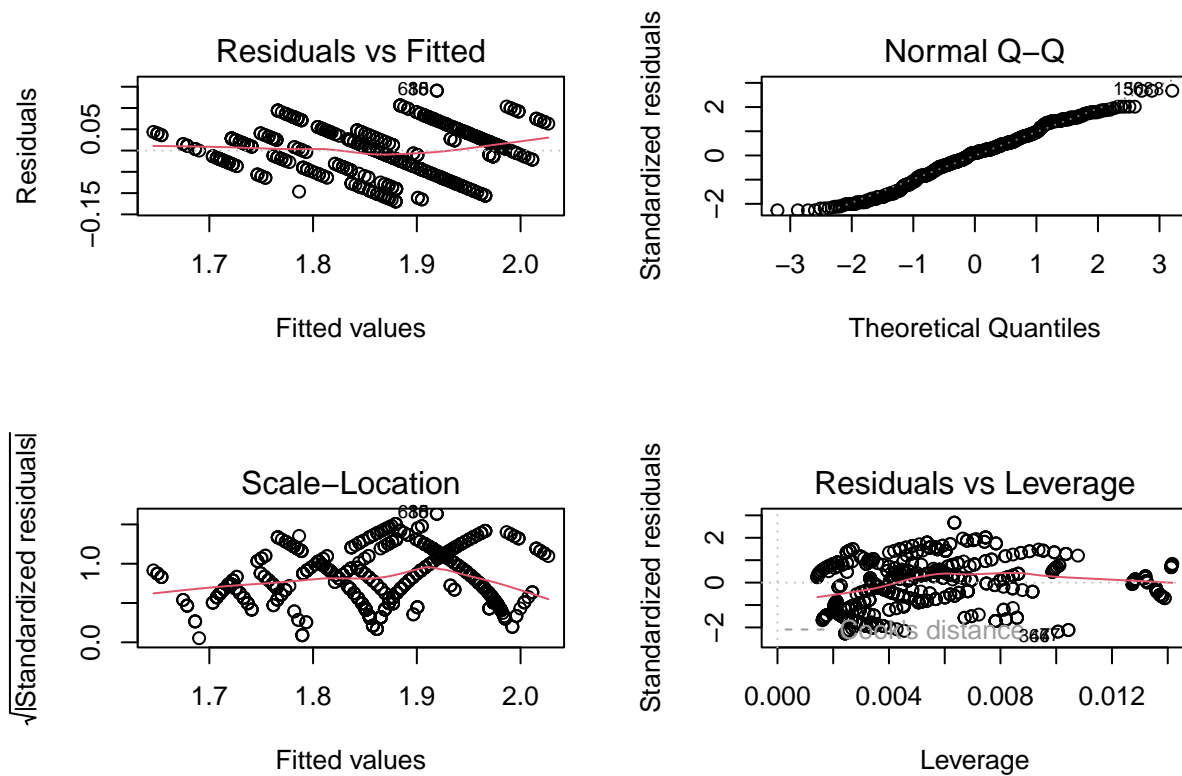
- Observations in the train dataset:  $n = 747$
- Explanatory variables:  $p = 3$
- Residual error degrees of freedom:  $DF = n - p - 1 = 747 - 3 - 1 = 743$

### 8.1.1 Diagnostic plots

```

par(mfrow = c(2,2))
plot(m)

```



```
par(mfrow = c(1,1))
```

### 8.1.2 Test the model with anova()

Compare the  $\Pr(>F)$  with the p-values from the summary of the model. They should be comparable.

```
anova(m)
```

```
## Analysis of Variance Table
##
## Response: PriceCH
##
##           Df Sum Sq Mean Sq F value    Pr(>F)
## PriceMM      1  2.8610  2.86102 1025.83 < 2.2e-16 ***
## WeekofPurchase 1  1.4268  1.42685  511.60 < 2.2e-16 ***
## STORE         1  1.2217  1.22168  438.04 < 2.2e-16 ***
## Residuals    743  2.0722  0.00279
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

### 8.1.3 Predict on test dataset




```
p <- predict(m, test)

# For ggplot we need a dataframe:
df <- data.frame(p, test)
```

#### 8.1.4 Plot predictions vs actuals

```
df %>% ggplot(aes(x = PriceCH, y = p)) +
  geom_point() +
  geom_smooth(method = 'lm', col = 'red', se=FALSE) +
  scale_y_continuous('Predictions') +
  scale_x_continuous('Price of Citrus Hill OJ (USD)' ) +
  ggtitle('Price of CH OJ', 'Dataset source: OJ{ISLR}')
```

```
## 'geom_smooth()' using formula = 'y ~ x'
```



#### 8.1.5 Assess performance: RMSE and $R^2$

- Root Mean Squared Error
- R-squared

```
# RMSE
sqrt(mean((test$PriceCH - p)^2))
```

```
## [1] 0.05600214
```

```
# R squared
cor(test$PriceCH, p)^2 ## R-Squared
```

```
## [1] 0.7140079
```

## 9 Example 4: Without Intercept

### 9.1 Fit the model without Intercept: Train set

Iterate to fit the model. See `anova()` description a few chunks below.

```
# You can start with a '+' on the first variable if you want to.
# I like to start with a '+' because it can flexibly add and delete variables
# as I go through the analysis process.
```

```
m_noi <- lm(PriceCH ~ -1
            +PriceMM
            +WeekofPurchase
            +STORE
            , train)
```

```
# Display the summary with the correlation of the coefficients.
summary(m_noi, corr = TRUE)
```

```
##
## Call:
## lm(formula = PriceCH ~ -1 + PriceMM + WeekofPurchase + STORE,
##     data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.133496 -0.034634  0.001591  0.049256  0.114761
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## PriceMM          0.2735939  0.0183128   14.94  <2e-16 ***
## WeekofPurchase  0.0049025  0.0001488   32.95  <2e-16 ***
## STORE           0.0293412  0.0015266   19.22  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.05903 on 744 degrees of freedom
## Multiple R-squared:  0.999, Adjusted R-squared:  0.999
## F-statistic: 2.497e+05 on 3 and 744 DF, p-value: < 2.2e-16
##
## Correlation of Coefficients:
```

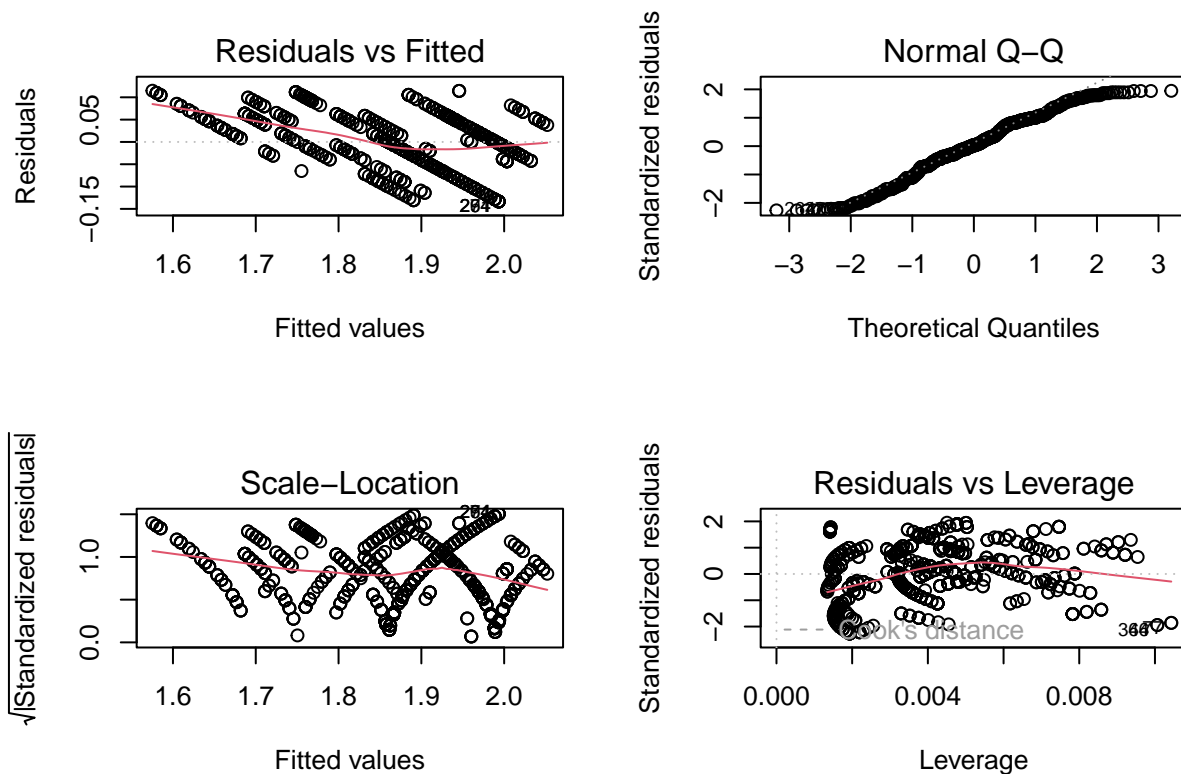
```
##           PriceMM WeekofPurchase
## WeekofPurchase -1.00
## STORE          -0.18    0.12
```

Note: Degrees of freedom =  $n - p - 1$ , when using the *Intercept* as we did above.

- Observations in the train dataset:  $n = 747$
- Explanatory variables:  $p = 3$
- Intercept = NO, think of it as = 0
- Residual error degrees of freedom:  $DF = n - p - 0 = 747 - 3 - 0 = 744$

### 9.1.1 Plots

```
par(mfrow = c(2,2))
plot(m_noi)
```



```
par(mfrow = c(1,1))
```

### 9.1.2 Predict on test dataset

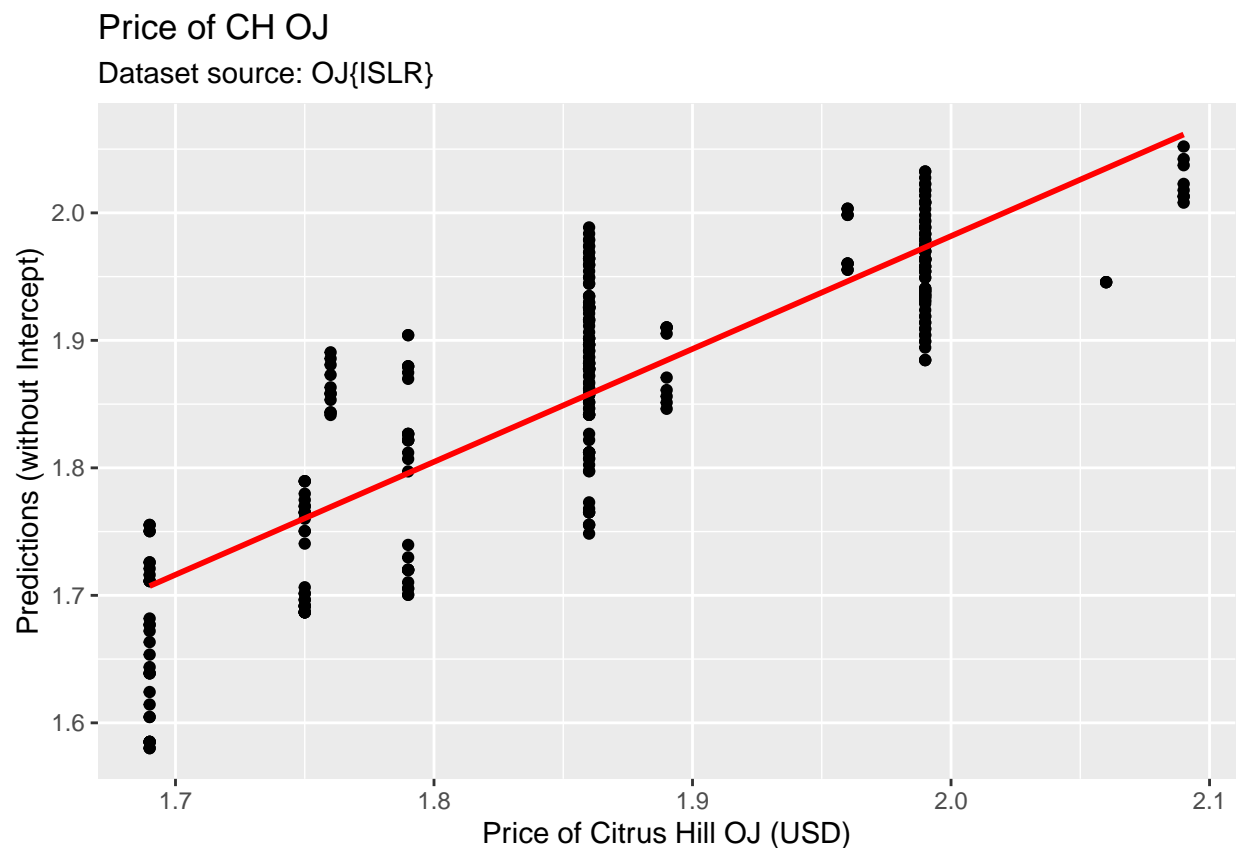
```
p_noi <- predict(m_noi, test)

# For ggplot we need a dataframe:
df_noi <- data.frame(p_noi, test)
```

### 9.1.3 Plot predictions vs actuals

```
df_noi %>% ggplot(aes(x = PriceCH, y = p_noi)) +
  geom_point() +
  geom_smooth(method = 'lm', col = 'red', se=FALSE) +
  scale_y_continuous('Predictions (without Intercept)') +
  scale_x_continuous('Price of Citrus Hill OJ (USD)') +
  ggtitle('Price of CH OJ', 'Dataset source: OJ{ISLR}')
```

```
## 'geom_smooth()' using formula = 'y ~ x'
```



### 9.1.4 Assess performance: RMSE and $R^2$

- Root Mean Squared Error
- R-squared

```
# RMSE
sqrt(mean((test$PriceCH - p_noi)^2))
```

```
## [1] 0.06115975
```

```
# R squared
cor(test$PriceCH, p_noi)^2 ## R-Squared
```

```
## [1] 0.7050249
```

Here the performance went down by a little bit only.

## 10 References

1. Harvard STAT 109 2023. Weekly slides by Dr. Bharatendra Rai.
2. Dr. Bharatendra Rai.YouTube channel. [https://youtu.be/cW59Yh\\_GfNk](https://youtu.be/cW59Yh_GfNk)
3. John Maindonald and W. John Braun. “Data Analysis and Graphics Using R”. Cambridge. Third Ed. ISBN 978-0-521-76293-9. 5th printing 2016.
4. Gareth James, et al. “And Introduction to Statistical Learning with Applications in R.” Springer Science. ISBN 978-1-4614-7137-0. 8th printing 2017.