

# Simulate continuous variables dataset

## Contents

<b>Purpose</b>	<b>1</b>
<b>Dependent variables</b>	<b>1</b>
Simulate the data . . . . .	1
Plot histograms and scatterplot . . . . .	2
<b>Independent variables</b>	<b>3</b>
Simulate the data . . . . .	3
Plot histograms and scatterplot . . . . .	3
<b>Simulate known linear regression</b>	<b>5</b>
Simulate the data . . . . .	5
Scatter plots . . . . .	5
Apply linear regression model . . . . .	6
Predict . . . . .	6
<b>References</b>	<b>8</b>

## Purpose

This script shows how to simulate a dataset that can be used in regression problems.

In regression, the variables that we are measuring are continuous, and the variable that we are predicting is continuous as well.

## Dependent variables

### Simulate the data

Create two Normally distributed datasets that have a relationship.

Play with the number of samples and we move the means around.

```

# From Harvard data science class (see references at the end of this notebook)
x <- rnorm(10000, mean=10, sd=sqrt(5))

# Initialize y with x...We would have a straight line if plotting y~x
y <- x

# Now inject variability to each, and we will not have a straight line exactly
x <- x + rnorm(10000, sd=2)
y <- y + rnorm(10000, sd=2)

```

## Plot histograms and scatterplot

```

br<- -5:25 # set manually bins for histograms
# save histograms for X and Y , don't plot yet
hx <- hist(x, breaks=br, plot=F)
hy <- hist(y, breaks=br, plot=F)


# prepare 2 panels in one plot:
old.par <- par(mfrow=c(1,2))

# plot histograms side by side using rbind
barplot(rbind(hx$density,hy$density),
        beside=T,
        col=c(rgb(0,0.2,1), rgb(0,1,0.3)),
        legend=c('X','Y'),
        main='Empirical distributions of X and Y',
        names=br[-1])


# Scatter plot
plot(x,y,
     xlab='X values',
     ylab='Y values',
     main='X vs Y scatterplot',
     pch=19,
     cex=0.3)

```

## Empirical distributions of X and



## X vs Y scatterplot



```
# restore graphical attributes to previous values:  
par(old.par)
```

## Independent variables

### Simulate the data

Create two independent Normally distributed datasets x and y.

Play with the number of samples and we move the means around.

```
# From Harvard data science class (see references at the end of this notebook)  
# simulate sampling of 10000 values for X and for Y.  
# We can play with the mean and sd. Should have same size to keep it balanced.  
x <- rnorm(10000, mean=10, sd=3)  
y <- rnorm(10000, mean=10, sd=3)
```

### Plot histograms and scatterplot

```
# Set manually bins for histograms  
br<- -5:25  
# Save histograms for X and Y , don't plot yet
```

```

hx <- hist(x, breaks=br, plot=F)
hy <- hist(y, breaks=br, plot=F)


# prepare 2 panels in one plot:
old.par <- par(mfrow=c(1,2))

# plot histograms side by side using rbind
barplot(rbind(hx$density,hy$density),
        beside=T,
        col=c(rgb(0,0.2,1), rgb(0,1,0.3)),
        legend=c('X','Y'),
        main='Empirical distributions of X and Y',
        names=br[-1])


# Scatter plot
plot(x,y,
     xlab='X values',
     ylab='Y values',
     main='X vs Y scatterplot',
     pch=19,
     cex=0.3)

```

**Empirical distributions of X and**



**X vs Y scatterplot**



```

# restore graphical attributes to previous values
par(old.par)

```

## Simulate known linear regression

In this approach, we will simulate data where we know the linear regression parameters.

### Simulate the data


Here we simulate X to be Uniformly distributed across a set of values.

We simulate Y with a known intercept, plus a slope times X with a random variability.

```
# X has a uniform distribution over a sequence over a range  
x <- seq(-3, 3, length=100)  
  
# Y is based on X with a slope, an intercept and normal randomness  
y <- 1 + 2*x+rnorm(100, sd=1)
```

### Scatter plots

```
plot(x, y, pch=19, cex=0.7)
```



## Apply linear regression model

```
m <- lm(y~x)
summary(m)

##
## Call:
## lm(formula = y ~ x)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.15705 -0.69116  0.06019  0.69113  2.42504
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.99052    0.09591   10.33  <2e-16 ***
## x            1.93098    0.05482   35.22  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.9591 on 98 degrees of freedom
## Multiple R-squared:  0.9268, Adjusted R-squared:  0.926
## F-statistic: 1241 on 1 and 98 DF,  p-value: < 2.2e-16
```


Notice the intercept is 1.01, while our empirical value was 1.

And the slope is 1.96, while our empirical value was 2.

## Predict

```
yp <- predict(m)

# Plot the prediction. Replot and add prediction points:
plot(x, y, pch=19, cex=0.7)
points(x,yp,col="red",pch=19,cex=0.7)
```



## Prediction metrix

```
rss <- sum((y - yp)^2)           # Sum of Squares Estimated (aka SSE)
rse <- sqrt(rss/98)
tss <- sum((y - mean(y))^2)       # Sum of Squares Total (aka SST)
ssr <- sum((yp - mean(y))^2)      # Sum of Squares Regression
se <- rse/sqrt(sum((x-mean(x))^2)) # Standard Error
roh_squared <- ssr / tss

cat(' RSS aka SSE = ', rss)
```

## RSS aka SSE = 90.1385

```
cat('\n RSE aka ?? = ', rse)
```

##

## RSE aka ?? = 0.9590519

```
cat('\n TSS aka SST = ', tss)
```

##

## TSS aka SST = 1231.342

```
cat('\n SSR = ', ssr)
```

```
##
```

```
## SSR = 1141.204
```

```
cat('\n Coefficient of determination (roh squared) = ', roh_squared)
```

```
##
```

```
## Coefficient of determination (roh squared) = 0.9267965
```

```
cat('\n roh = ', sqrt(roh_squared))
```

```
##
```

```
## roh = 0.9627027
```

```
#
```

```
# cat('\n SE = ', se)
```

```
# (tss-rss)/tss
```

## References

- Harvard “Elements of Statistical Learning” (2021) taught by professors Dr. Sivachenko, Dr. Farutin